

# RED-BLACK TREES

- *Red-black trees:*

- Binary search trees augmented with node color
- Operations designed to guarantee that the height  $h = O(\lg n)$

- First: describe the properties of red-black trees
- Then: prove that these guarantee  $h = O(\lg n)$
- Finally: describe operations on red-black trees

## RED-BLACK PROPERTIES

- The *red-black properties*:

1. Every node is either red or black
2. Every leaf (NULL pointer) is black
  - Note: this means every “real” node has 2 children
3. If a node is red, both children are black
  - Note: can't have 2 consecutive reds on a path
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

## RED-BLACK TREES

- Put example on board and verify properties:
  1. Every node is either red or black
  2. Every leaf (NULL pointer) is black
  3. If a node is red, both children are black
  4. Every path from node to descendent leaf contains the same number of black nodes
  5. The root is always black
- *black-height*: # black nodes on path to leaf
  - Label example with  $h$  and  $bh$  values

## HEIGHT OF RED-BLACK TREES

- *What is the minimum black-height of a node with height  $h$ ?*
- A: a height- $h$  node has black-height  $\geq h/2$
- Theorem: A red-black tree with  $n$  internal nodes has height  $h \leq 2 \lg(n + 1)$
- *How do you suppose we'll prove this?*

## RB TREES: PROVING HEIGHT BOUND

- Prove:  $n$ -node RB tree has height  $h \leq 2 \lg(n+1)$
- Claim: A subtree rooted at a node  $x$  contains at least  $2^{bh(x)} - 1$  internal nodes
  - Proof by induction on height  $h$
  - Base step:  $x$  has height 0 (i.e., NULL leaf node)
    - *What is  $bh(x)$ ?*

## RB TREES: PROVING HEIGHT BOUND

- Prove:  $n$ -node RB tree has height  $h \leq 2 \lg(n+1)$
- Claim: A subtree rooted at a node  $x$  contains at least  $2^{bh(x)} - 1$  internal nodes
  - Proof by induction on height  $h$
  - Base step:  $x$  has height 0 (i.e., NULL leaf node)
    - *What is  $bh(x)$ ?*
    - A: 0
    - So...subtree contains  $2^{bh(x)} - 1$   
 $= 2^0 - 1$   
 $= 0$  internal nodes (TRUE)

## RB TREES: PROVING HEIGHT BOUND

- Inductive proof that subtree at node  $x$  contains at least  $2^{\text{bh}(x)} - 1$  internal nodes
  - Inductive step:  $x$  has positive height and 2 children
    - Each child has black-height of  $\text{bh}(x)$  or  $\text{bh}(x)-1$  (*Why?*)
    - The height of a child = (height of  $x$ ) - 1
    - So the subtrees rooted at each child contain at least  $2^{\text{bh}(x)-1} - 1$  internal nodes
    - Thus subtree at  $x$  contains
$$(2^{\text{bh}(x)-1} - 1) + (2^{\text{bh}(x)-1} - 1) + 1$$
$$= 2 \cdot 2^{\text{bh}(x)-1} - 1 = 2^{\text{bh}(x)} - 1 \text{ nodes}$$

## RB TREES: PROVING HEIGHT BOUND

- Thus at the root of the red-black tree:

$$n \geq 2^{\text{bh}(\text{root})} - 1 \quad (\text{Why?})$$

$$n \geq 2^{h/2} - 1 \quad (\text{Why?})$$

$$\lg(n+1) \geq h/2 \quad (\text{Why?})$$

$$h \leq 2 \lg(n+1) \quad (\text{Why?})$$

Thus  $h = O(\lg n)$




## RB TREES: WORST-CASE TIME

- So we've proved that a red-black tree has  $O(\lg n)$  height
- Corollary: These operations take  $O(\lg n)$  time:
  - Minimum(), Maximum()
  - Successor(), Predecessor()
  - Search()
- Insert() and Delete():
  - Will also take  $O(\lg n)$  time
  - But will need special care since they modify tree

## RED-BLACK TREES: AN EXAMPLE

○ *Color this tree:*




Red-black properties:

1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# RED-BLACK TREES: THE PROBLEM WITH INSERTION

○ *Insert 8:*

■ *Where does it go?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# RED-BLACK TREES: THE PROBLEM WITH INSERTION

## ○ Insert 8

- *Where does it go?*
- *What color should it be?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# RED-BLACK TREES: THE PROBLEM WITH INSERTION

## ○ Insert 8

- *Where does it go?*
- *What color should it be?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# RED-BLACK TREES: THE PROBLEM WITH INSERTION

## ○ Insert 11

- *Where does it go?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# RED-BLACK TREES: THE PROBLEM WITH INSERTION

## ○ Insert 11

- *Where does it go?*
- *What color?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# RED-BLACK TREES: THE PROBLEM WITH INSERTION

## ○ Insert 11

- *Where does it go?*
- *What color?*
  - Can't be red! (#3)




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black



# RED-BLACK TREES: THE PROBLEM WITH INSERTION

## ○ Insert 11

- *Where does it go?*
- *What color?*
  - Can't be red! (#3)
  - Can't be black! (#4)




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# RED-BLACK TREES: THE PROBLEM WITH INSERTION

## ○ Insert 11

- *Where does it go?*
- *What color?*
- Solution:  
recolor the tree




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# RED-BLACK TREES: THE PROBLEM WITH INSERTION

## ○ Insert 10

- *Where does it go?*

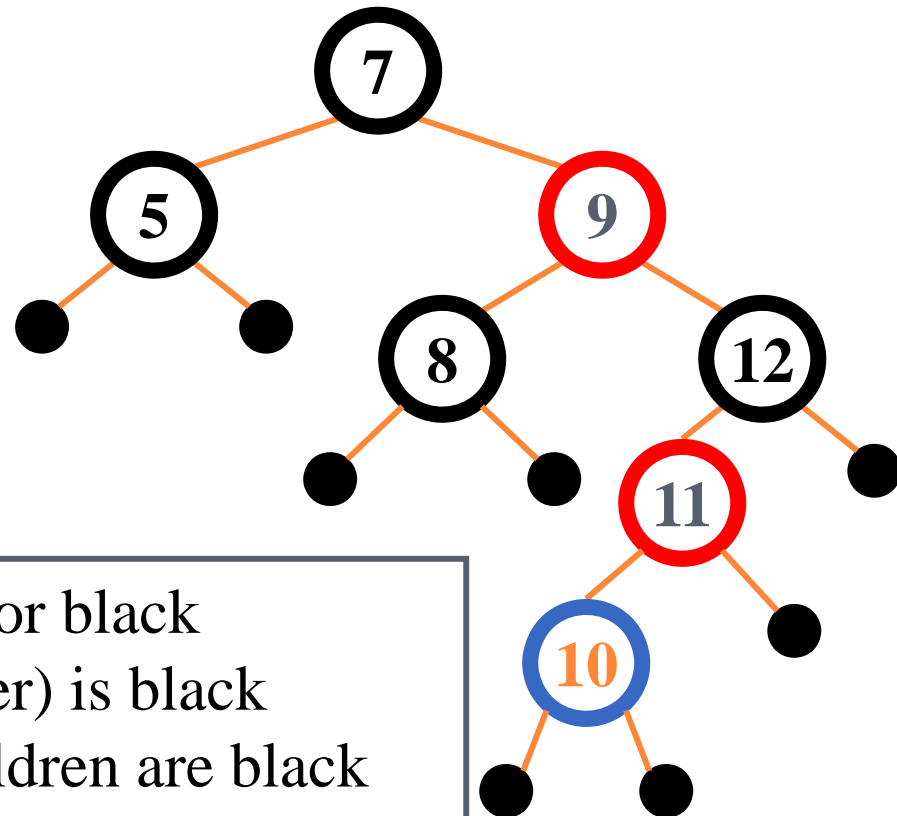


1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# RED-BLACK TREES: THE PROBLEM WITH INSERTION

## ○ Insert 10

- *Where does it go?*
- *What color?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# RED-BLACK TREES: THE PROBLEM WITH INSERTION


## ○ Insert 10

- *Where does it go?*
- *What color?*
  - A: no color! Tree is too imbalanced
  - Must change tree structure to allow recoloring
- Goal: restructure tree in  $O(\lg n)$  time




## RB TREES: ROTATION

- Our basic operation for changing tree structure is called *rotation*:



- *Does rotation preserve inorder key ordering?*
- *What would the code for **rightRotate()** actually do?*


## RB TREES: ROTATION



- Answer: A lot of pointer manipulation
  - $x$  keeps its left child
  - $y$  keeps its right child
  - $x$ 's right child becomes  $y$ 's left child
  - $x$ 's and  $y$ 's parents change
- *What is the running time?*

## ROTATION EXAMPLE


- Rotate left about 9:





## ROTATION EXAMPLE

- Rotate left about 9:



# RED-BLACK TREES: INSERTION

## ○ Insertion: the basic idea

- Insert  $x$  into tree, color  $x$  red
- Only r-b property 3 might be violated (if  $p[x]$  red)
  - If so, move violation up tree until a place is found where it can be fixed
- Total time will be  $O(\lg n)$

```
rbInsert(x)
```

```
    treeInsert(x);
```

```
    x->color = RED;
```

```
    // Move violation of #3 up tree, maintaining #4 as invariant:
```

```
    while (x!=root && x->p->color == RED)
```

```
    if (x->p == x->p->p->left)
```

```
        y = x->p->p->right;
```

```
        if (y->color == RED)
```

```
            x->p->color = BLACK;
```

```
            y->color = BLACK;
```

```
            x->p->p->color = RED;
```

```
            x = x->p->p;
```

```
        else    // y->color == BLACK
```

```
            if (x == x->p->right)
```

```
                x = x->p;
```

```
                leftRotate(x);
```

```
            x->p->color = BLACK;
```

```
            x->p->p->color = RED;
```

```
            rightRotate(x->p->p);
```

```
    else    // x->p == x->p->p->right
```

```
        (same as above, but with
```

```
        "right" & "left" exchanged)
```

} Case 1

} Case 2

} Case 3

```
rbInsert(x)
```

```
    treeInsert(x);
```

```
    x->color = RED;
```

```
    // Move violation of #3 up tree, maintaining #4 as invariant:
```

```
    while (x!=root && x->p->color == RED)
```

```
    if (x->p == x->p->p->left)
```

```
        y = x->p->p->right;
```

```
        if (y->color == RED)
```

```
            x->p->color = BLACK;
```

```
            y->color = BLACK;
```

```
            x->p->p->color = RED;
```

```
            x = x->p->p;
```

```
        else // y->color == BLACK
```

```
            if (x == x->p->right)
```

```
                x = x->p;
```

```
                leftRotate(x);
```

```
                x->p->color = BLACK;
```

```
                x->p->p->color = RED;
```

```
                rightRotate(x->p->p);
```

```
    else // x->p == x->p->p->right
```

```
        (same as above, but with
```

```
        "right" & "left" exchanged)
```

} Case 1: uncle is RED

} Case 2

} Case 3

## RB INSERT: CASE 1


if (y->color == RED)

  x->p->color = BLACK;


  y->color = BLACK;

  x->p->p->color = RED;

  x = x->p->p;



case 1  
.....>



Change colors of some nodes, preserving #4: all downward paths have equal b.h.  
The while loop now continues with x's grandparent as the new x

## RB INSERT: CASE 1


if (y->color == RED)

  x->p->color = BLACK;

  y->color = BLACK;

  x->p->p->color = RED;


  x = x->p->p;




Same action whether x is a left or a right child

## RB INSERT: CASE 2

```
if (x == x->p->right)
    x = x->p;
    leftRotate(x);
// continue with case 3 code
```



case 2  
.....>




Transform case 2 into case 3 (x is left child) with a left rotation  
This preserves property 4: all downward paths contain same number of black nodes

## RB INSERT: CASE 3

**x->p->color = BLACK;**

**x->p->p->color = RED;**

**rightRotate(x->p->p);**



- Case 3:
  - “Uncle” is black
  - Node  $x$  is a left child
- Change colors; rotate right

Perform some color changes and do a right rotation

Again, preserves property 4: all downward paths contain same number of black nodes



## RB INSERT: CASES 4-6

- Cases 1-3 hold if  $x$ 's parent is a left child
- If  $x$ 's parent is a right child, cases 4-6 are symmetric (swap left for right)